

Journal Name

Crossmark

PAPER

RECEIVED
dd Month yyyy

REVISED
dd Month yyyy

Quantum Kernels for EEG: An Invariance Mismatch Masquerading as Advantage

Md. Intiaj Alam Sajin^{1,*} , Md Wahiduzzaman Suva¹, Esm E Moula Chowdhury Abha¹,
Samiul Islam Saif¹  and Muhammad Hasibur Rashid Chayon¹ 

¹Department of Computer Science, American International University-Bangladesh (AIUB), Dhaka, Bangladesh

*Author to whom any correspondence should be addressed.

E-mail: 26-94090-2@student.aiub.edu, imtiajsajin@gmail.com (M I A Sajin); 26-94088-2@student.aiub.edu (M W Suva); 26-94089-2@student.aiub.edu (E E M C Abha); Md.samiulislam.aiub@gmail.com (S I Saif); chayon@aiub.edu (M H R Chayon)

Abstract

Objective. Quantum machine learning is often proposed for electroencephalography (EEG) decoding, usually without controls showing that the quantum part does the work. We ask whether quantum kernels beat well-tuned classical baselines, and if not, why.

Approach. A trace-normalised EEG covariance matrix is a quantum density matrix, so quantum information geometry applies directly. We benchmark Hilbert–Schmidt, fidelity, Bures and quantum-relative-entropy kernels against tuned Riemannian, common spatial pattern and filter-bank baselines on three motor-imagery datasets under one nested cross-validation budget. The quantum quantities are invariant only under orthogonal congruence, while EEG’s nuisances span the full congruence group that classical baselines handle. Borrowing recentring from Riemannian transfer learning, we refer each state to a training-set Fréchet mean, prove this restores exact affine invariance, and add a classical twin differing only in metric.

Main results. In the sensor frame the classical baselines won on every dataset. Recentring raised Gram variance several-fold and improved every kernel, by up to +0.102 on PhysioNet and +0.187 on IV-2a, reversing the headline comparison from +0.0635 for the best classical pipeline to −0.0120 against it. The twin shows this is not quantum. Across seven settings (three datasets, three register sizes, within-subject, cross-subject and cross-session evaluation) two one-sided tests place the two families within ± 0.028 accuracy in all 35 comparisons. 3 of those intervals exclude zero, 2 favouring a quantum kernel and 1 the twin; every one stays inside the margin, and no quantum kernel beats the best classical pipeline in any setting.

Significance. We find no quantum advantage and trace the apparent one to an invariance mismatch rather than kernel concentration: supply the missing invariance and the quantum kernels match their classical counterparts. Recentring is a precondition for a fair comparison; omitting it is a route to a false positive.

Keywords: quantum machine learning, brain–computer interface, electroencephalography, quantum kernels, Riemannian geometry, motor imagery, benchmarking

1 Introduction

Decoding motor imagery from EEG is a core problem in brain–computer interface (BCI) research, and the field has converged on a small set of strong classical methods: common spatial patterns (CSP) [1] and, more recently, classifiers built on the Riemannian geometry of spatial covariance matrices [2, 3]. Large-scale benchmarking finds that covariance tangent-space projection and CSP consistently define the strongest methodological families across public motor-imagery datasets; Lotte *et al* [4] and Yger *et al* [5] review the field and the Riemannian branch of it respectively.

In parallel, a fast-growing literature applies quantum machine learning (QML) to EEG, typically reporting improvements over a classical comparator. It now spans motor imagery [6, 7, 8, 9], biosignal encoding [10], drowsiness detection [11], clinical classification [12] and, most recently, a hybrid quantum-classical decoder evaluated on real quantum hardware in a clinical journal [13]; a systematic review of QML across digital health appeared in 2025 [14].

These reports share a methodological pattern that makes them difficult to interpret, and two recent examples show it is not a matter of care but of missing controls. QEEGNet [6] appends a variational circuit to EEGNet and reports a mean test accuracy of 0.381 against EEGNet’s 0.377 on four-class IV-2a: an improvement, but over one classical architecture, at an absolute level far below what filter-bank and Riemannian pipelines reach on that dataset, so the comparison cannot establish that the quantum layer helps relative to the state of the art. Carter *et al* [13] combine a spatial-temporal deep network with digitised counterdiabatic quantum features and reach 0.898 accuracy on binary motor imagery across 51 participants, which is genuinely strong; and they report, to their credit, that the quantum features change the classification of only 244 of 30,590 trials and that the improvement over their own classical model does not reach significance overall ($p = 0.073$ by McNemar’s test), reaching it only in a younger subgroup. Both results are consistent with a quantum component that is doing little, and neither study is designed to tell the difference.

This is not a complaint that the field is careless; it is a structural problem with the question. Schuld and Killoran [15] argue that chasing quantum advantage as a headline is the wrong goal for quantum machine learning precisely because it licenses comparisons the experiment cannot support, and the recommendation this paper ends with is the constructive form of that argument: report what the quantum structure contributes against a comparator that differs from it in one respect only.

The pattern behind this is threefold. Classical baselines are often weak or untuned: plain support vector machines on band power rather than the Riemannian or filter-bank CSP methods that define the state of the art. Hyperparameters are frequently optimised for the quantum model alone. Most critically, two controls are almost universally absent: an *entanglement ablation*, which tests whether the quantum resource contributes anything at all, and a *dimension-matched classical comparator*, without which an apparent quantum effect is confounded with the dimensionality reduction applied before encoding. This paper adds a third control, a comparator that differs from the quantum kernel in metric alone, and it is that one that turns out to be decisive.

Quantum kernels are the natural object of study here. Supervised quantum models with fixed encodings are kernel methods [16, 17], so a kernel comparison captures the whole model class rather than one architecture, and it sidesteps the trainability problems that afflict variational circuits [18, 19]. What is known about their usefulness is mixed. There is a constructed problem with a proven exponential separation [20], but on natural data the advantage is bounded by how much the quantum kernel’s inductive bias happens to match the target [21, 22], and general surveys [23] have not translated into demonstrated wins on biological signals.

The wider QML literature gives strong reasons for caution. Bowles *et al* [24] benchmarked twelve quantum models across 160 datasets and found that out-of-the-box classical models generally outperformed them; more pointedly, removing entanglement often left performance equal or better. Thanasilp *et al* [25] proved that quantum kernel values concentrate exponentially in qubit count toward a fixed value, collapsing the model to a trivial one. Neither result has been systematically tested on EEG.

This paper makes four contributions.

First, we identify a *principled* quantum encoding for EEG. An EEG trial’s spatial covariance matrix is symmetric positive definite; normalised to unit trace it is a bona fide quantum density matrix (figure 1). This is an identity, not an analogy, and it makes quantum information geometry (Uhlmann fidelity, the Bures metric, von Neumann entropy) applicable to EEG natively, with no lossy squeezing of features into rotation angles. The Hilbert–Schmidt overlap $\text{tr}(\rho\sigma)$ is moreover exactly the quantity a SWAP test estimates, so the construction maps onto a hardware-native primitive. Related tooling exists [26], but its quantum step operates on flattened tangent-space vectors, discarding the density-matrix structure.

Second, we report a controlled benchmark of 23 pipelines on 104 subjects with nested cross-validation, matched tuning budgets, both of the missing controls, and paired subject-wise statistics with multiple-comparison correction, replicated on IV-2a and extended to filter-bank baselines, cross-subject and cross-session transfer, and a register sweep to six qubits in both frames.

Third, and centrally, we show that the comparison such benchmarks perform, ours included, in its first form, is confounded by the coordinate frame. The quantum quantities above are invariant under orthogonal congruence only, while EEG’s nuisance transformations (electrode gain, reference choice, source mixing, and the change from one subject to the next) generate the full congruence group, under which every strong classical baseline *is* invariant. We prove that referring each state to a Fréchet-mean reference state restores exact affine invariance while remaining a legitimate quantum operation, and we show that this single change relieves concentration by $4.7\text{--}9.2\times$ and reverses the sign of the classical-versus-quantum comparison on both datasets.

Fourth, we show that the reversal is nonetheless *not* evidence of quantum advantage, using a control we believe should be standard in this literature: a classical SPD-manifold kernel placed in the same classifier, on the same data, in the same frame, differing from the quantum kernels only in the metric. No quantum kernel separates from that twin by more than 0.028 in any of seven settings, including cross-session transfer, the regime in which the shift our invariance argument does not cover is proportionally largest, and none beats the best classical pipeline. The practical implication is uncomfortable but concrete: a quantum kernel compared against a classical baseline in a different frame will show an advantage of roughly the size we report, produced entirely by the frame.




Figure 1. From scalp EEG to a quantum state, computed from real recordings rather than simulated. (a) Band-passed sensorimotor traces for one left-hand and one right-hand trial. (b) Scalp topography of mu/beta power, right minus left: power falls over the hemisphere contralateral to the imagined hand, which is the physiological effect every decoder in this study exploits. (c) The spatial covariance of a single trial over the eight decoding channels. (d) The same matrix divided by its trace. It is now a density matrix, and (e) its eigenvalues are non-negative and sum to one, so they are a population distribution. The step from (c) to (d) is the whole encoding: no dimensionality reduction, no rotation-angle embedding.

We state our scope plainly. At the channel counts considered here (8–64, i.e. 3–6 qubits) every kernel we evaluate is classically computable in $O(n_{\text{ch}}^3)$. We make no claim of quantum speedup; the object of study is whether quantum *geometry* improves decoding. Every quantum kernel is evaluated by exact state-vector simulation: there is no gate error, no decoherence and no readout bias, and section 3.11 adds only sampling noise. Hardware would make the quantum side worse, never better, so the comparison throughout is maximally generous to it and a negative result here is not weakened by the idealisation.

The contribution we intend for neural engineering is practical rather than quantum. Covariance decoders are the field’s working standard, and the two devices this paper turns on the quantum kernels, recentring to a label-free reference state and a comparator that differs in metric alone, apply to any new covariance kernel a practitioner might be offered, quantum or otherwise. The invariance proposition says *why* recentring is a precondition for comparing geometries at all, where the Riemannian transfer-learning literature [27, 28] has so far justified it empirically; and the negative result is stated with enough control that it can be built on rather than merely cited.

2 Methods

2.1 Data and preprocessing

We use the PhysioNet EEG Motor Movement/Imagery Database (EEGMMIDB) [29, 30], recorded with the BCI2000 system at 160 Hz over 64 channels. We analyse the standard binary protocol: imagined left-hand versus right-hand movement (runs 4, 8 and 12), giving 45 trials per subject. Signals are band-pass filtered to 8–30 Hz, spanning the μ and β rhythms in which motor-imagery




Figure 2. Study design, and where the control sits. Arrows give the order of processing: EEG to State to Frame to Kernels to Evaluation. **EEG**: three public motor-imagery datasets, PhysioNet EEGMMIDB (104 subjects, 45 trials), BCI Competition IV-2a (9 subjects, 288 trials, two sessions) and Cho2017 (52 subjects), reduced to the same eight sensorimotor electrodes (filled), the same 8–30 Hz band and the same one-second window, so every difference reported here is a property of the method and not of the preprocessing. **State**: each trial becomes an 8×8 spatial covariance C ; divided by its trace it is a density matrix ρ on $\log_2 8 = 3$ qubits, with no dimensionality reduction and no angle embedding. **Frame**: EEG’s nuisances (electrode gain, head geometry, skull conductivity, source mixing) act by congruence $C \mapsto ACA^T$. The density-matrix kernels are invariant only under the orthogonal subgroup in the sensor frame, and exactly invariant under the full group in the reference frame of (6), where $W = M^{-1/2}$ and the Fréchet mean M is estimated label-free from training trials alone. Every suite is run in both frames. **Kernels**: four pipeline families under one protocol. The bracket marks the paper’s central control: the classical Riemannian twin and the quantum kernels consume the same covariances, sit in the same support vector machine, receive the same tuning budget and are expressed in the same frame, so the only thing that differs between them is the metric. **Evaluation**: nested cross-validation (5 folds $\times$ 3 repeats outer, 4-fold inner, identical budget for every pipeline), paired per-subject Wilcoxon signed-rank with Holm correction inside each pre-specified family, and two one-sided tests for equivalence against the twin, across the five evaluation settings. The editable source of this figure is distributed with the manuscript as `architecture.drawio`.

event-related desynchronisation appears, and epoched from 0.5 to 3.5 s after cue onset to exclude the cue-evoked transient. Epochs are resampled to 128 Hz.

The primary analysis uses eight sensorimotor channels (FC3, FCz, FC4, C3, Cz, C4, CP3, CP4), a register of exactly three qubits under the amplitude encoding described below. Subjects 88, 89, 92, 100 and 104 are excluded for documented annotation and sampling-rate defects. We analyse the first 104 eligible subjects.

Spatial covariance matrices are estimated per trial with the oracle approximating shrinkage (OAS) estimator, which is well conditioned at the trial counts available here.

2.2 Second dataset: BCI Competition IV-2a

A single dataset cannot separate a property of the methods from a property of the data. PhysioNet gives breadth but only 45 trials per subject, so a negative result there invites the objection that the quantum models were starved. We therefore repeat the entire benchmark on BCI Competition IV-2a [31] (9 subjects, 288 trials each, 2025 fold scores), accessed through MOABB [32]. The two datasets sit at opposite corners of the design space: many subjects with little data each versus few subjects with a great deal.

IV-2a is also the dataset the quantum-EEG literature reports on [6], so results here are directly comparable to published claims. The same eight sensorimotor channels are used, all of which exist in the IV-2a montage, so the register size is identical at three qubits. Band-pass, resampling, pipelines, hyperparameter grids and cross-validation protocol are unchanged; only the data differs. Each dataset keeps its own standard epoch window, since each competition defines its own cue timing, and the comparison of interest is whether the *ranking* of pipelines moves rather than whether absolute accuracies coincide.

2.3 Third dataset: Cho2017

The twin comparison is a claim of sameness, and sameness is bounded by the number of subjects. PhysioNet and IV-2a together give 104 and 9 paired observations; the Cho *et al* motor-imagery dataset [33], accessed through MOABB, adds 52 subjects recorded on a 64-channel Biosemi ActiveTwo system at 512 Hz, with 200 imagined left- and right-hand movements per subject (240 for 3 of them) in a single session. It is the largest public left-versus-right motor-imagery corpus and was chosen for that reason alone: it tightens the equivalence bounds of section 3.5, and it supplies a third estimate of the frame effect at a trial count between the other two. The same eight sensorimotor channels, band-pass, resampling, pipelines, grids and cross-validation protocol are used, with the dataset's own cue window.

2.4 EEG covariance matrices as quantum states

Let $C_i \in \mathbb{R}^{n_{\text{ch}} \times n_{\text{ch}}}$ be the spatial covariance of trial i . Since C_i is symmetric positive definite,

$$\rho_i = \frac{C_i}{\text{tr } C_i} \quad (1)$$

satisfies $\rho_i \succeq 0$ and $\text{tr } \rho_i = 1$: it is a density matrix on an n_{ch} -dimensional Hilbert space, that is, on $\log_2 n_{\text{ch}}$ qubits. Sweeping the channel count over powers of two therefore sweeps the register size with no change of method, a property we exploit in section 3.3.

2.5 Density-matrix kernels

Equation (1) makes the following quantum-information quantities directly available as kernels between trials.

The *Hilbert–Schmidt overlap*

$$K_{\text{HS}}(\rho, \sigma) = \text{tr}(\rho\sigma) \quad (2)$$

is an inner product on Hermitian matrices, so its Gram matrix is positive semi-definite by construction and requires no repair. It is precisely the observable estimated by a SWAP test.

The *Uhlmann fidelity* [34]

$$F(\rho, \sigma) = \left(\text{tr} \sqrt{\sqrt{\rho} \sigma \sqrt{\rho}} \right)^2 \quad (3)$$

is the quantum generalisation of the Bhattacharyya coefficient. We evaluate it via the eigenvalues of the symmetrised product, which is numerically stable for the matrix sizes involved.

Both quantities proved severely concentrated on EEG (section 3.3), so we additionally define bandwidth-parameterised kernels that retain the quantum geometry while restoring a length scale, exactly as a radial basis function (RBF) kernel does for the Euclidean distance:

$$K_{\text{HS-RBF}}(\rho, \sigma) = \exp(-\gamma \operatorname{tr}[(\rho - \sigma)^2]), \quad (4)$$

$$K_{\text{Bures}}(\rho, \sigma) = \exp(-\gamma d_B^2(\rho, \sigma)), \quad d_B^2 = 2 \left(1 - \sqrt{F(\rho, \sigma)}\right). \quad (5)$$

Here d_B is the Bures distance [35], the geodesic distance of the Bures metric on state space, which for unit-trace matrices coincides with the Bures–Wasserstein distance occasionally used in the Riemannian EEG literature as a distance but, to our knowledge, never as a kernel. The bandwidth γ is a bandwidth, and bandwidth is known to control quantum-kernel expressivity and to be decisive for whether such kernels generalise at all [36]. It is set by the median heuristic computed *on the training split only*, with a multiplier selected in the inner cross-validation loop.

2.6 Reference states, and the invariance the sensor frame lacks

Equation (1) expresses each state in the frame the electrodes happen to define. That frame is not neutral.

The nuisance transformations of EEG act on the spatial covariance by congruence, $C \mapsto ACA^\top$ with A invertible: per-electrode gain and impedance, the choice of reference electrode, linear source mixing and volume conduction, and the change from one subject or session to the next all take this form. The affine-invariant Riemannian metric is invariant under the whole congruence group, which is the structural reason tangent-space decoding is the classical state of the art and why it transfers.

The quantum quantities are not. $\operatorname{tr}(\rho\sigma)$, the Uhlmann fidelity, the Bures distance and the quantum relative entropy are invariant only under the *orthogonal* subgroup $A \in O(n_{\text{ch}})$. This is an asymmetry in the comparison itself, not merely a property of the data: pyRiemann’s tangent-space estimator fits a reference mean M and maps $C \mapsto \log(M^{-1/2}CM^{-1/2})$, and both the minimum-distance-to-mean classifier and CSP are congruence-invariant by construction. Every strong classical baseline in our suite therefore possesses an invariance that none of the sensor-frame quantum kernels possess.

The repair is not new, and we want to be plain about that. Whitening covariances by their Fréchet mean is *Riemannian recentring*, introduced by Zanini *et al* [27] to align sessions and subjects and extended by Rodrigues *et al* [28] into Riemannian Procrustes Analysis; He and Wu [37] give a Euclidean counterpart. In the classical setting it is understood as a domain-adaptation step, applied because it helps in transfer, and it is optional: the affine-invariant methods it is applied to are already congruence-invariant without it.

Our point is that for quantum kernels the same operation plays a different and stronger role. It is not an optional aid to transfer but the step that supplies an invariance the kernels structurally lack, and which the baselines they are being compared against already have. We therefore refer each state to a *reference state* M , the Fréchet mean of the *training* covariances, via the whitening $W = M^{-1/2}$:

$$\tilde{\rho}_i = \frac{WC_iW}{\operatorname{tr}(WC_iW)}. \quad (6)$$

Proposition. *Under a common congruence $C_i \mapsto AC_iA^\top$ with A invertible, the kernels $\operatorname{tr}(\tilde{\rho}\tilde{\sigma})$, $F(\tilde{\rho}, \tilde{\sigma})$, $d_B(\tilde{\rho}, \tilde{\sigma})$ and $J(\tilde{\rho}, \tilde{\sigma})$ are invariant.*

Proof. The affine-invariant mean is equivariant, $M \mapsto AMA^\top$. Put $B = (AMA^\top)^{-1/2}A$. Then $BMB^\top = I = WMW^\top$, so $U := BM^{1/2}$ satisfies $UU^\top = I$ and $B = UW$. Hence $WC_iW \mapsto U(WC_iW)U^\top$ for one common orthogonal U , and trace normalisation commutes with orthogonal conjugation. Each quantity above is invariant under a common unitary. $\square$

In the reference frame the quantum kernels are thus *exactly* affine-invariant, sharing the invariance group of the affine-invariant Riemannian metric. Figure 3 sets out the geometry and verifies the invariance numerically. Two consequences follow, and we test both. First, the comparison against the classical baselines becomes like for like. Second, and less obviously, the quantum kernels can now at best *match* the Riemannian metric on any nuisance that is a congruence; parity, not superiority, is the ceiling the algebra permits.

The proposition’s scope should be stated exactly, because the transfer experiments go beyond it. It covers a congruence *common* to every state being compared, which is the situation within one subject and one session. In cross-subject and cross-session evaluation we whiten each subject or session by its own mean. If subject or session s is distorted by its own A_s , the same argument turns each distortion into an orthogonal U_s , but the U_s differ between subjects or sessions, and nothing in the whitening removes that relative rotation. Neither the quantum kernels nor the

Riemannian twin is invariant to it, so the proposition does not predict parity there. Whether parity survives that residual, together with any change in spectral shape or conditioning, is an empirical question, and sections 3.7 and 3.8 test it.

To state the contribution precisely: the alignment is borrowed, the reason for needing it is not. What we add is the observation that quantum kernels lack an invariance the classical baselines already possess, a proof that recentring supplies exactly that invariance for the quantum quantities, and the control in section 3.5 showing that once it is supplied nothing quantum remains to distinguish them.

Equation (6) is not classical preprocessing smuggled into a quantum method. The map $\rho \mapsto W\rho W/\text{tr}(W\rho W)$ is a filtering (Lüders) operation with Kraus operator W followed by renormalisation, so it is state preparation and the SWAP-test implementation path is unchanged. M is estimated on training data only.




Figure 3. (a) EEG covariances cluster tightly about their Fréchet mean M , and a nuisance congruence moves the whole cluster. Whitening by M spreads the states about the identity and reduces the nuisance to a common rotation, to which every quantum invariant is blind. (b) The proposition verified numerically: reference-frame residuals sit at machine precision, sensor-frame ones are $O(10^{-1})$. (c) Off-diagonal Gram variance rises by 4.7–9.2 $\times$ in the reference frame. (d) The sensor-frame overlaps saturate near unity in every subject ($n = 104$).

We also add the symmetrised quantum relative entropy

$$J(\rho, \sigma) = \frac{1}{2} \text{tr}[(\rho - \sigma)(\log \rho - \log \sigma)], \quad (7)$$

used as $\exp(-\gamma J)$. It is distinct from the classical Kullback–Leibler divergence between the Gaussians the covariances parameterise, which pyRiemann supplies as a distance and which involves $\text{tr}(\sigma^{-1}\rho)$ and log-determinants rather than matrix logarithms.

2.7 Circuit embedding kernels and the entanglement ablation

For comparison with the prevailing approach, we also evaluate parameterised-circuit embedding kernels $K(x, z) = |\langle \phi(x) | \phi(z) \rangle|^2$ (figure 4), where ϕ is an instantaneous-quantum-polynomial (IQP) style feature map of the Havlíček *et al* family [38], built from Hadamard and R_Z rotations with data-dependent ZZ couplings. Inputs are Riemannian tangent-space features reduced by PCA to the number of qubits and scaled to a rotation-angle range; the angle scale is itself a tuned hyperparameter. A variant with fixed CNOT entanglers is also included.

The **entanglement ablation** is the decisive control: an otherwise identical circuit with every entangling gate deleted. The resulting kernel factorises into single-qubit kernels and is classically trivial. If it matches or exceeds the entangled version, quantum structure is not contributing [24].

Kernels are evaluated by preparing each embedded state once on a state-vector simulator, which is mathematically identical to running the adjoint kernel-estimation circuit at infinite shots but costs $O(n2^q)$ rather than $O(n^2)$ circuit executions.




Figure 4. The circuit feature maps, at three qubits with one layer shown. Rotation angles carry the PCA-reduced tangent-space features and the kernel is the state overlap $|\langle\phi(x)|\phi(z)\rangle|^2$. The IQP map uses data-dependent ZZ couplings; the ring map uses fixed CNOTs. The third panel is the entanglement ablation: the identical circuit with every entangling gate deleted, which factorises the kernel into single-qubit terms and isolates what entanglement contributes. The figure is rendered by Qiskit from a circuit that is checked against the executed PennyLane circuit before drawing: the two agree on full Gram matrices to $< 10^{-15}$, so the diagram cannot drift from the code it documents.

2.8 Classical baselines and controls

Five classical baselines represent the genuine state of the art rather than a strawman: log-variance with shrinkage LDA; CSP with shrinkage LDA; minimum distance to Riemannian mean (MDM); Riemannian tangent space with logistic regression; and tangent space with an RBF SVM.

Four controls isolate specific confounds: the entanglement ablation described above; *dimension-matched* linear and RBF SVMs operating on exactly the same PCA features the circuit kernels receive; and a log-Euclidean tangent-space classifier as a first classical geometric comparator for the density-matrix kernels.

Together with the two circuit kernels and the five sensor-frame density-matrix kernels (Hilbert–Schmidt overlap, fidelity, HS-RBF, Bures-RBF and QRE-RBF), these make up the sixteen pipelines of the *core suite*, all operating in the sensor frame. The *extended suite* adds seven: the same five density-matrix kernels in the reference frame (6), and the two classical SPD kernels described next, for twenty-three pipelines in all. Table S1 lists every hyperparameter grid and every fixed setting. Where a comparison names “the best quantum kernel” without qualification it means the most accurate quantum pipeline of the suite in question; the density-matrix comparisons of section 3.5 always name the kernel explicitly.

The control the argument ultimately rests on is stricter than that, and we define it here because everything in section 3.5 depends on it. A *metric-matched twin* is a classical kernel on the SPD manifold used exactly as the reference-frame density-matrix kernels are used: the same covariances, referred to the same training-set Fréchet mean, a precomputed Gram matrix in the same support vector machine, and the same grid over the SVM cost C . Only the metric differs. Our twin is the Riemannian kernel of pyRiemann [39], $k_R(X, Y) = \text{tr}[\log(M^{-1/2}XM^{-1/2})\log(M^{-1/2}YM^{-1/2})]$ with M the training-set Fréchet mean, the inner product of the two states in the tangent space at the reference. It has no bandwidth, so it is tuned over C alone, as are the two overlap kernels; the three bandwidth-parameterised quantum kernels additionally tune a bandwidth multiplier. Because a common congruence maps $M^{-1/2}XM^{-1/2}$ to $U(M^{-1/2}XM^{-1/2})U^T$ for an orthogonal U , exactly as in the proposition, k_R has the same affine invariance as the reference-frame quantum kernels, which is why it is the twin in every setting. We also report pyRiemann’s log-Euclidean kernel, $\text{tr}[(\log X - \log M)(\log Y - \log M)]$, as a second classical kernel. It centres in the log domain rather than whitening, is not congruence-invariant, and so is not a twin; it is shown because it is sometimes the stronger classical kernel. Because the twin shares every element of the quantum pipeline except the metric, a difference between them is attributable to the metric.

Three further suites reuse this design. A *filter-bank* suite splits the 8–30 Hz band into four sub-bands (8–12, 12–16, 16–22 and 22–30 Hz) and stacks them, so that one covariance spans $4 \times 8 = 32$ dimensions, a single density matrix on five qubits, and adds filter-bank CSP [40] and filter-bank tangent-space baselines alongside filter-bank versions of every kernel and twin. A *cross-subject* suite pools the trials of all but one subject, trains on the pool, and tests on the held-out subject; its hyperparameters are chosen by subject-grouped inner cross-validation on the training subjects, so no held-out subject influences any choice. A *cross-session* suite uses the two

recording sessions of IV-2a: for each subject, and in both directions, it trains on one session and tests on the other, choosing hyperparameters by stratified inner cross-validation within the training session alone. In its reference frame each session is whitened by its own Fréchet mean, which uses no labels and is, for the test session, exactly the unsupervised adaptation a deployed decoder can perform on the unlabelled data it receives before decoding anything. All three suites are described where their results appear.

2.9 Evaluation protocol and statistics

Each subject is evaluated independently. The outer loop is a repeated stratified k -fold (5 folds $\times$ 3 repeats) estimating generalisation; the inner loop is a 4-fold grid search selecting hyperparameters. Crucially, *every* pipeline, including all classical baselines, receives the same inner search budget. All fitted transformations (tangent-space reference mean, scalars, PCA, and the kernel bandwidth) are refit within each training split, so no information leaks across folds.

Across subjects we follow the standard recommendation for comparing classifiers over multiple datasets [41] and use the Wilcoxon signed-rank test on paired per-subject accuracies, with Holm correction [42] across the family of comparisons, and report the paired effect size d_z . Per-subject optimality varies substantially in motor-imagery decoding, so paired non-parametric testing (rather than comparison of grand means) is the appropriate inference.

Two points about the correction structure, since several families appear below. Holm correction is applied *within* each family, and the families are the pre-specified ones: the suite-wide comparison against the reference pipeline, the pre-specified contrasts of table 4, and the metric-matched twin comparisons. These were fixed by the design of the study, one family per question asked, and not selected after inspecting results; no comparison reported here was chosen because it was significant. We deliberately do not apply a single global correction across every family, because the questions are logically distinct and a global correction would make the null results, which are the paper’s actual claims, easier rather than harder to obtain.

The outer partition is itself a source of variation, so the PhysioNet three-qubit suite, on which the decisive comparison is first made, is repeated in full under two further outer-partition seeds (table 8); nothing else changes between repetitions.

Where the claim is that two methods are the *same*, a non-significant difference is not sufficient evidence, so we additionally report two one-sided tests (TOST) on the paired differences against a margin of ± 0.02 accuracy. The margin was not pre-registered: it was chosen after the first within-subject PhysioNet and IV-2a analyses, the filter-bank analysis and the cross-subject analysis, and before every other analysis, and it was not moved afterwards. Because of that, table 7 also reports the smallest margin at which each equivalence would hold, so the result can be read at any margin (table 7). The margin, its justification and the achieved equivalence bounds are given there.

2.10 Implementation

Analyses use MNE-Python [43], pyRiemann, scikit-learn and PennyLane [44]. Code and results are available at <https://github.com/Intiaj-Sajin/quantaEEG>.

3 Results

3.1 Quantum kernels do not outperform tuned classical baselines

Table 1 reports within-subject performance for all 16 pipelines; Figure 5 shows the same data with per-subject spread. Quantum kernels occupy five of the bottom six positions. Absolute accuracies are modest because EEGMMIDB contains many near-chance subjects, which is precisely why paired testing is required. They are also in line with the field: the MOABB reproducibility benchmark [45] reports 65.8–68.2% for CSP and tangent-space pipelines on this dataset and paradigm using the full 64-channel montage, against which our 0.629 on eight channels is the expected reduction, not a sign of a mistuned baseline.

The primary pre-specified comparison, best classical versus best quantum, is significant: classical/CSP+LDA exceeded quantum/CNOT-kernel-SVM by $\Delta = +0.027$ ($p < 0.001$, $d_z = 0.34$), and was superior in 70/104 subjects (Table 4).

Against the Riemannian tangent-space reference (Table 2, Figure S1), 8 of 15 comparisons survived Holm correction. 9 reached $p < 0.05$ uncorrected, and 8 of those favour the reference rather than the pipeline tested against it. No quantum kernel exceeded the reference at any significance level: the best of them falls short by -0.0153 .

Table 1. Within-subject decoding performance on PhysioNet EEGMMIDB left- versus right-hand motor imagery. Accuracy and AUC are means over subjects of the per-subject nested cross-validation score; SD is between subjects. Runtime is mean wall-clock seconds per subject for the full hyperparameter search. Rows are ordered by accuracy.

Pipeline	Group	Accuracy	SD	AUC	Runtime (s)
CSP + LDA	Classical	0.629	0.165	0.664	3.98
TS + LR	Classical	0.618	0.174	0.651	4.95
TS + RBF-SVM	Classical	0.617	0.166	0.644	13.10
PCA-matched-linear	Control	0.606	0.164	0.627	5.10
MDM	Classical	0.603	0.164	0.636	0.48
CNOT-kernel-SVM	Quantum	0.602	0.156	0.626	19.06
PCA-matched-RBF	Control	0.601	0.159	0.621	14.16
IQP-no-entangle	Control	0.601	0.156	0.627	18.09
logeuclid-TS + LR	Control	0.595	0.149	0.629	4.46
IQP-kernel-SVM	Quantum	0.581	0.136	0.604	21.45
logvar + LDA	Classical	0.574	0.131	0.609	0.09
Bures-RBF-SVM	Quantum	0.566	0.128	0.586	14.47
QRE-RBF-SVM	Quantum	0.565	0.127	0.586	12.84
HS-RBF-SVM	Quantum	0.553	0.117	0.571	10.90
Fidelity-SVM	Quantum	0.542	0.112	0.583	5.41
HS-overlap-SVM	Quantum	0.531	0.098	0.579	3.90

Table 2. Paired comparisons against the TS + LR reference across subjects (Wilcoxon signed-rank, two-sided). Δ is the mean per-subject accuracy difference; d_z is the paired effect size; p_{Holm} is corrected across the whole family of comparisons. 8 of 15 comparisons survive correction.

Pipeline	Δ accuracy	p	p_{Holm}	d_z	Better in
CSP + LDA	+0.0116	0.048	0.334	+0.168	53/104
TS + RBF-SVM	-0.0006	0.520	0.814	-0.011	47/104
PCA-matched-linear	-0.0122	0.407	0.814	-0.144	50/104
MDM	-0.0144	0.192	0.767	-0.189	46/104
CNOT-kernel-SVM	-0.0153	0.140	0.698	-0.182	44/104
PCA-matched-RBF	-0.0168	0.116	0.696	-0.196	39/104
IQP-no-entangle	-0.0168	0.194	0.767	-0.192	45/104
logeuclid-TS + LR	-0.0229	0.005	0.040	-0.271	40/104
IQP-kernel-SVM	-0.0368	<0.001	0.003	-0.381	33/104
logvar + LDA	-0.0443	<0.001	0.007	-0.362	40/104
Bures-RBF-SVM	-0.0519	<0.001	0.005	-0.390	36/104
QRE-RBF-SVM	-0.0530	<0.001	0.003	-0.394	36/104
HS-RBF-SVM	-0.0650	<0.001	<0.001	-0.457	34/104
Fidelity-SVM	-0.0755	<0.001	<0.001	-0.518	34/104
HS-overlap-SVM	-0.0870	<0.001	<0.001	-0.578	29/104




Figure 5. Within-subject accuracy by pipeline ($n = 104$ subjects, nested cross-validation). Bars are means over subjects; grey dots are individual subject means. The dashed line marks chance. Every pipeline, classical baselines included, received the same inner hyperparameter search.

3.2 The controls: entanglement and dimension matching

Both control ablations point in the direction predicted by [24], but neither reaches significance on accuracy (Table 4). Removing entanglement changed accuracy by $+0.0200$ ($p < 0.001$, superior in 68/104 subjects); the dimension-matched linear SVM exceeded the IQP circuit kernel by $+0.0246$ ($p < 0.001$).

We regard the honest reading as follows: deleting the quantum resource did not harm performance and numerically helped, but accuracy alone, measured on a noisy and low-SNR signal, does not have the power to establish the effect. The next section shows that the same effect is unambiguous when measured on the kernel itself.

3.3 Kernel concentration, and why entanglement makes it worse

Raw overlap kernels are severely concentrated on EEG. Off-diagonal Gram entries had mean 0.982 ± 0.018 (Hilbert–Schmidt) and 0.981 ± 0.016 (fidelity): every pair of trials appears almost identical, the Gram matrix is effectively rank-one, and the classifier has nothing to separate. This is the pathology proved by Thanasilp *et al* [25], here reproduced on real biological data rather than synthetic constructions.

Exploiting (1), we swept register size by varying the channel count over powers of two (4/8/16/32/64 channels = 2–6 qubits) with no change of method (Table S2, Figure 6). Two findings emerge.

First, **entanglement accelerates the collapse**. The entangled IQP kernel loses variance by a factor of 0.424 per qubit, more than halving with each added qubit, while the identical circuit with entanglers deleted decays at 0.727. In log-slope the entangled kernel concentrates $2.7\times$ faster (-0.859 versus -0.319 per qubit), a $30\times$ versus $3.5\times$ loss across the range. This supplies a mechanism for the observation of [24] that removing entanglement often improves QML models: entanglement drives embedded states toward mutual orthogonality, pushing the Gram matrix toward the identity, a memorising, non-generalising model.

Second, **the two kernel families concentrate for different reasons**. The raw density-matrix kernels do not follow the qubit-count law at all; their variance *increases* with dimension (factors 1.45–1.52). They are already saturated at two qubits because EEG covariance matrices are intrinsically similar to one another, a property of the data, not of the register. The practical corollary inverts standard practice: the density-matrix route should use *more* channels, not fewer.




Figure 6. Kernel variance versus register size on real EEG ($n = 104$ subjects). Lower is worse: a kernel whose variance collapses cannot separate trials. Line style encodes the kernel family. The entangled circuit kernel (orange, dashed) decays fastest.

Table 3. Replication on BCI Competition IV-2a (9 subjects, 288 trials each), with the PhysioNet result repeated for comparison. Protocol, channels, tuning budget and pipelines are identical; only the data differs. Quantum kernels occupy the four lowest positions on both datasets.

Pipeline	Group	Acc (IV-2a)	AUC (IV-2a)	Acc (Phys.)
TS + LR	Classical	0.763	0.822	0.618
TS + RBF-SVM	Classical	0.755	0.811	0.617
CSP + LDA	Classical	0.751	0.808	0.629
logeuclid-TS + LR	Control	0.751	0.810	0.595
MDM	Classical	0.723	0.777	0.603
logvar + LDA	Classical	0.697	0.748	0.574
PCA-matched-linear	Control	0.690	0.735	0.606
PCA-matched-RBF	Control	0.687	0.729	0.601
CNOT-kernel-SVM	Quantum	0.685	0.726	0.602
IQP-no-entangle	Control	0.682	0.722	0.601
IQP-kernel-SVM	Quantum	0.672	0.712	0.581
Bures-RBF-SVM	Quantum	0.600	0.630	0.566
HS-RBF-SVM	Quantum	0.595	0.621	0.553
HS-overlap-SVM	Quantum	0.586	0.629	0.531
Fidelity-SVM	Quantum	0.580	0.616	0.542

3.4 Replication on IV-2a: more data widens the gap

Repeating the entire benchmark on BCI Competition IV-2a reproduces the ordering and sharpens it (Table 3, Figure 7). The best classical pipeline reached 0.763 against 0.685 for the best quantum kernel and 0.580 for the worst, a classical-versus-quantum gap of 0.079 where PhysioNet gave 0.027, and a gap to the weakest quantum kernel of 0.183 against 0.099.

The direction of that change is the point. Classical methods gained roughly 0.15 accuracy from the six-fold increase in trials; the density-matrix kernels gained about 0.05. **More and cleaner data made the quantum kernels relatively worse, not better**, which is the opposite of what a data-starvation explanation predicts, and exactly what section 3.3 implies: a kernel whose off-diagonal entries all sit near unity gains no additional structure from additional trials, while the classical methods gain a great deal.

The ranking is also stable. Spearman rank correlation between the two orderings is 0.851 across all 16 pipelines, and the four lowest positions are occupied by the same four quantum kernels on both datasets with zero rank shift. This is a property of the methods, not of the dataset.

Table 4. Pre-specified comparisons on both datasets, combined by Fisher’s method. Positive Δ favours the first-named pipeline. At $n = 9$ the two-sided signed-rank test cannot return $p < 0.0039$, so the IV-2a entries marked $\dagger$ sit at that floor: the test is saturated rather than merely significant. Both classical-versus-quantum contrasts hold in every one of the 9 IV-2a subjects.

Comparison	Dataset	Δ acc	p	Better
Best classical vs. best quantum	PhysioNet	+0.0269	<0.001	70/104
	IV-2a	+0.0662	0.004 $\dagger$	9/9
	<i>Fisher combined</i>		<0.001	
Entanglement ablation	PhysioNet	+0.0200	<0.001	68/104
	IV-2a	+0.0100	0.203	7/9
	<i>Fisher combined</i>		<0.001	
Dimension-matched control	PhysioNet	+0.0246	<0.001	67/104
	IV-2a	+0.0178	0.129	7/9
	<i>Fisher combined</i>		<0.001	

Statistically (Table 4), on IV-2a the best classical pipeline beat the best quantum kernel by +0.066 ($p = 0.004$, $d_z = 1.14$) in 9/9 subjects. CSP+LDA, the classical comparator fixed on PhysioNet, beat the sensor-frame fidelity kernel by +0.171 ($d_z = 1.45$) in 9/9 subjects. Combining the two independent datasets by Fisher’s method gives $p < 0.001$ and $p < 0.001$ respectively.

One caveat must be stated precisely rather than left to the reader. With n pairs the two-sided signed-rank test cannot return a p below 2^{1-n} ; at $n = 9$ that floor is 0.0039, so a Holm correction over 15 comparisons cannot fall below 0.0547 whatever the effect size. Ten of the fourteen IV-2a family comparisons sit exactly at that floor. Reporting them as “not significant after correction” would misdescribe a saturated test as a null result, which is why the pre-specified comparisons in Table 4, which require no family correction, carry the inference.




Figure 7. The same 16 pipelines on two datasets at opposite corners of the design space. Identical channels, protocol, tuning budget and code; only the data differs. Line slope is the quantity of interest: classical and control pipelines rise steeply with the extra trials while the four density-matrix and overlap kernels stay low. The dashed line marks chance.

3.5 The sensor frame is a confound, and correcting it reverses the result

Every result above evaluates the density-matrix kernels in the sensor frame. Section 2.6 argued that this is not neutral: the quantum kernels are invariant only under orthogonal congruence, while the tangent-space, MDM and CSP baselines they are being compared against are invariant under the full congruence group that EEG’s nuisances generate. We now re-run the entire benchmark with the density-matrix kernels expressed relative to a training-set reference state (6), changing nothing else. The core pipelines reproduce their values in table 1 to ≤ 0.0001 , so the two suites are directly comparable.

Three consequences follow, and the third is the one that matters.

First, concentration is largely an artefact of the frame. Measured over 104 subjects, the off-diagonal variance of the Gram matrix rises by a factor of 4.7–9.2 on moving to the reference frame. The sensor-frame Hilbert–Schmidt kernel has off-diagonal mean 0.982 and standard deviation 0.023; a nuisance congruence of the kind an electrode-gain change produces shifts individual kernel entries by an order of magnitude more than that entire discriminative spread. The kernel was measuring the frame, not the trial.

The register sweep of section 3.3 was made in the sensor frame, and its second finding, that the density-matrix kernels are saturated by the data rather than by the register, has to be re-examined now that most of that saturation is known to be an artefact of the frame. Repeating the sweep in both frames (table S3, figure S2(b)) gives a cleaner answer. In the reference frame no kernel’s Gram variance falls below its three-qubit value at any larger register, and at six qubits it stands at 2.0–3.8 times that value; the exponential collapse with qubit count of [25] is absent from this family in either frame. The gain from recentring narrows from 4.7–9.2 $\times$ at three qubits to 1.7–3.2 $\times$ at six, but that is because the sensor-frame variance itself rises 5.7–10.7 $\times$ over the same range, not because the reference frame deteriorates. The corollary of section 3.3, that the density-matrix route tolerates wider registers, therefore survives the correction, with one qualification: the reference-frame variance is not strictly monotone, three of the four kernels dipping by up to 0% between five and six qubits.

Second, accuracy improves substantially and consistently (table 5). On PhysioNet every kernel improves, by +0.064 to +0.102 ($n = 104$, all $p \leq 0.001$). On IV-2a the same corrections are worth

Table 5. Effect of referring the density-matrix kernels to a reference state. Each kernel is evaluated twice under an identical protocol, differing only in whether the states are expressed in the sensor frame (1) or relative to the training-set Fréchet mean (6). Δ is the mean per-subject accuracy gain from the reference frame, tested by paired Wilcoxon signed-rank across subjects; p_{Holm} corrects for every row of this table as one family. Every kernel improves on every dataset.

Dataset	Kernel	Sensor	Reference	Δ	p	p_{Holm}	Better
PhysioNet, 3 q ($n = 104$)	Fidelity	0.542	0.641	+0.0991	<0.001	<0.001	71/104
	HS overlap	0.531	0.633	+0.1021	<0.001	<0.001	75/104
	HS-RBF	0.553	0.629	+0.0761	<0.001	<0.001	69/104
	Bures-RBF	0.566	0.631	+0.0650	<0.001	<0.001	70/104
	QRE-RBF	0.565	0.629	+0.0644	<0.001	<0.001	71/104
PhysioNet, 4 q ($n = 104$)	Fidelity	0.536	0.627	+0.0912	<0.001	<0.001	68/104
	HS overlap	0.526	0.630	+0.1033	<0.001	<0.001	72/104
	HS-RBF	0.550	0.625	+0.0745	<0.001	<0.001	69/104
	Bures-RBF	0.566	0.632	+0.0663	<0.001	<0.001	69/104
	QRE-RBF	0.562	0.631	+0.0689	<0.001	<0.001	71/104
Cho2017, 3 q ($n = 52$)	Fidelity	0.615	0.648	+0.0332	<0.001	<0.001	40/52
	HS overlap	0.579	0.642	+0.0633	<0.001	<0.001	47/52
	HS-RBF	0.600	0.637	+0.0367	<0.001	<0.001	42/52
	Bures-RBF	0.627	0.639	+0.0117	0.009	0.023	33/52
	QRE-RBF	0.629	0.641	+0.0125	0.002	0.013	36/52
IV-2a, 3 q ($n = 9$)	Fidelity	0.580	0.767	+0.1873	0.004	0.023	9/9
	HS overlap	0.586	0.762	+0.1760	0.004	0.023	9/9
	HS-RBF	0.595	0.763	+0.1678	0.004	0.023	9/9
	Bures-RBF	0.600	0.765	+0.1646	0.004	0.023	9/9
	QRE-RBF	0.599	0.764	+0.1644	0.004	0.023	9/9

+0.164 to +0.187, and every kernel improves in *every* subject. That the effect is two to three times larger on IV-2a is expected rather than surprising: the reference state is a Fréchet mean estimated from training covariances, and 288 trials per subject estimate it far better than 45 do. Cho2017, with 200 trials per subject, shows that the estimate is not the only determinant: its correction is the smallest of the three, +0.012 to +0.063 ($n = 52$, all $p \leq 0.009$), because its sensor-frame kernels start much closer to the classical baselines (the best classical pipeline leads the best sensor-frame kernel by +0.0152 there, against +0.0635 on PhysioNet). The correction is worth most where the sensor frame is worst.

This also revises the interpretation of section 3.4. There, more data widened the classical-quantum gap, which we read as evidence that the quantum kernels could not exploit additional trials. The narrower and correct reading is that a kernel with the wrong invariance cannot exploit them. Given the right frame, the extra data is not merely usable but more valuable to the quantum kernels than to the classical baselines.

The headline comparison consequently reverses. Against classical/CSP+LDA, the best sensor-frame quantum kernel loses by +0.0635 ($p < 0.001$, classical better in 70/104 subjects); the best reference-frame kernel, quantum/Fidelity-ref-SVM, instead leads by 0.0120 ($p = 0.095$, classical better in only 44/104). This does not reach conventional significance, and we do not claim it as an effect. We dwell on it because it is the size and shape of result that this literature routinely reports as a positive finding, and because everything needed to present it as one (two datasets, a consistent direction, a larger margin on the better-powered data) is in hand at this point in the analysis.

Table 6. The decisive control. Each quantum kernel is compared against an SPD-manifold kernel used in the same support vector machine, on the same covariances, with the same tuning budget, in the same reference frame: only the metric differs. In every setting the twin is the affine-invariant Riemannian kernel, the one classical kernel here that shares the reference-frame quantum kernels' invariance under congruence. The log-Euclidean kernel lacks that invariance and is shown for comparison only. Columns give the best quantum kernel, the twin, the log-Euclidean kernel, the range of the five quantum-minus-twin differences, and the smallest p obtained by *any* of the five quantum kernels against the twin. Table 8 repeats the PhysioNet three-qubit row under two further partitions.

Setting	n	Quantum	Twin	Log-Eucl.	Δ (range over 5)	$\min p$
PhysioNet, 3 q	104	0.641	0.630	0.635	[−0.0012, +0.0113]	0.003
PhysioNet, 4 q	104	0.632	0.634	0.639	[−0.0088, −0.0017]	0.048
PhysioNet, 5 q, filter bank	104	0.640	0.639	0.637	[−0.0051, +0.0014]	0.362
PhysioNet, transfer (LOSO)	104	0.674	0.669	0.669	[+0.0028, +0.0053]	0.095
BCI IV-2a, 3 q	9	0.767	0.766	0.763	[−0.0033, +0.0016]	0.426
BCI IV-2a, cross-session	9	0.760	0.755	0.755	[−0.0085, +0.0054]	0.500
Cho2017, 3 q	52	0.648	0.638	0.636	[−0.0008, +0.0100]	0.001

Third, it is not a quantum advantage, and the control that establishes this is the central result of

this paper. A reference-frame quantum kernel differs from a classical Riemannian kernel in exactly one respect: the metric. Both consume the same covariances, both are used as precomputed kernels in the same support vector machine, both receive the same tuning budget, and after (6) both are affine-invariant. Table 6 reports that comparison in seven settings: three datasets, three register sizes, and within-subject, cross-subject and cross-session evaluation.

In none of them does a quantum kernel separate from its classical twin by more than 0.028 (figure 9), and in only one, discussed below, is a separation significant at the primary partition. On PhysioNet the differences span -0.0012 to +0.0113 with a smallest $p = 0.003$ in that setting; the twin itself, control/riemann-kernel-SVM, reaches 0.630 and is the second-best pipeline of the whole extended suite. The apparent quantum advantage is therefore attributable to two things, neither of them quantum: expressing the states in the reference frame, and using an SPD-manifold kernel in a support vector machine rather than a tangent-space projection.

Table 7. Equivalence of each quantum kernel to its classical twin, by two one-sided tests on the paired per-subject differences. A non-significant difference would only be an absence of evidence; TOST instead takes a *difference* as its null, so rejecting it supports equivalence. The margin $m = 0.02$ accuracy is not pre-registered. It was chosen after the first within-subject analyses of PhysioNet and IV-2a, the filter-bank analysis and the cross-subject analysis had been run, from the size of the reference-frame effects they showed, and before every other analysis reported here. It is left unchanged although some later effects are smaller (the correction now spans +0.012 to +0.217). CI is the 90 % interval, which corresponds to TOST at $\alpha = 0.05$; “bound” is the smallest margin at which equivalence would hold, so a reader preferring a stricter margin can read the answer off directly.

Setting	Kernel	n	Δ	90 % CI	p_{TOST}	Equiv.	Bound
PhysioNet, 3 qubits	Fidelity	104	+0.0113	[+0.0043, +0.0182]	0.0201	✓	0.0182
	HS overlap	104	+0.0027	[-0.0047, +0.0101]	0.0001	✓	0.0101
	HS-RBF	104	-0.0012	[-0.0096, +0.0072]	0.0002	✓	0.0096
	Bures-RBF	104	+0.0008	[-0.0068, +0.0083]	0.0000	✓	0.0083
	QRE-RBF	104	-0.0009	[-0.0088, +0.0069]	0.0001	✓	0.0088
PhysioNet, 4 qubits	Fidelity	104	-0.0067	[-0.0136, +0.0003]	0.0010	✓	0.0136
	HS overlap	104	-0.0038	[-0.0108, +0.0033]	0.0001	✓	0.0108
	HS-RBF	104	-0.0088	[-0.0159, -0.0016]	0.0054	✓	0.0159
	Bures-RBF	104	-0.0017	[-0.0077, +0.0043]	0.0000	✓	0.0077
	QRE-RBF	104	-0.0029	[-0.0087, +0.0029]	0.0000	✓	0.0087
PhysioNet, 5 qubits	Fidelity	104	+0.0014	[-0.0062, +0.0091]	0.0001	✓	0.0091
	HS overlap	104	-0.0051	[-0.0138, +0.0037]	0.0027	✓	0.0138
	HS-RBF	104	-0.0034	[-0.0117, +0.0048]	0.0006	✓	0.0117
	Bures-RBF	104	+0.0011	[-0.0060, +0.0081]	0.0000	✓	0.0081
	QRE-RBF	104	-0.0001	[-0.0074, +0.0072]	0.0000	✓	0.0074
Transfer (LOSO)	Fidelity	104	+0.0030	[-0.0016, +0.0076]	0.0000	✓	0.0076
	HS overlap	104	+0.0053	[-0.0004, +0.0111]	0.0000	✓	0.0111
	HS-RBF	104	+0.0034	[-0.0031, +0.0099]	0.0000	✓	0.0099
	Bures-RBF	104	+0.0043	[-0.0012, +0.0097]	0.0000	✓	0.0097
	QRE-RBF	104	+0.0028	[-0.0028, +0.0083]	0.0000	✓	0.0083
BCI IV-2a, 3 qubits	Fidelity	9	+0.0016	[-0.0049, +0.0080]	0.0004	✓	0.0080
	HS overlap	9	-0.0033	[-0.0109, +0.0043]	0.0018	✓	0.0109
	HS-RBF	9	-0.0028	[-0.0084, +0.0028]	0.0002	✓	0.0084
	Bures-RBF	9	-0.0008	[-0.0073, +0.0058]	0.0003	✓	0.0073
	QRE-RBF	9	-0.0019	[-0.0081, +0.0043]	0.0003	✓	0.0081
IV-2a cross-session	Fidelity	9	+0.0042	[-0.0057, +0.0142]	0.0094	✓	0.0142
	HS overlap	9	-0.0004	[-0.0124, +0.0117]	0.0082	✓	0.0124
	HS-RBF	9	-0.0085	[-0.0281, +0.0111]	0.1531	×	0.0281
	Bures-RBF	9	+0.0054	[-0.0138, +0.0246]	0.0975	×	0.0246
	QRE-RBF	9	+0.0054	[-0.0144, +0.0252]	0.1034	×	0.0252
Cho2017, 3 qubits	Fidelity	52	+0.0100	[+0.0047, +0.0153]	0.0013	✓	0.0153
	HS overlap	52	+0.0041	[-0.0010, +0.0093]	0.0000	✓	0.0093
	HS-RBF	52	-0.0008	[-0.0082, +0.0066]	0.0000	✓	0.0082
	Bures-RBF	52	+0.0010	[-0.0063, +0.0083]	0.0000	✓	0.0083
	QRE-RBF	52	+0.0034	[-0.0036, +0.0104]	0.0001	✓	0.0104

A non-significant difference is only an absence of evidence, and the claim we are making deserves better than that, so we test equivalence directly (table 7). Two one-sided tests on the paired differences, with a margin of ± 0.02 accuracy, establish equivalence in 32 of 35 comparisons; across all 35 the agreement is within ± 0.028 . The evidence is strongest where it should be, on the two better-powered datasets: on IV-2a and on Cho2017 all 5 and 5 comparisons are equivalent, with the widest intervals reaching only 0.0109 and 0.0153.

Two honest qualifications. First, the 2 comparisons that miss the ± 0.02 margin in the quantum kernel’s favour are balanced by 1 that miss it in the twin’s, which is the pattern expected of noise about zero rather than of a suppressed effect. Second, 3 of the 35 intervals exclude zero, mixed in

Table 8. The PhysioNet three-qubit comparison under three different outer cross-validation partitions (seeds). Each row repeats the full extended suite with new fold assignments; the inner search, the reference state and every other element are unchanged. Columns give the frame effect (range over the five kernels), the best reference-frame quantum kernel and the classical twin, the quantum-minus-twin difference (range over the five kernels), the smallest p against the twin, and how many of the five comparisons pass two one-sided tests at the pre-specified ± 0.02 margin, with the largest bound. The frame effect and the uncontrolled reversal reproduce in every partition; the twin comparison moves against the bandwidth-parameterised quantum kernels under the other two partitions and never in their favour (see text).

Seed	Frame Δ	Quantum	Twin	Quantum – twin	min p	TOST
0	[+0.064, +0.102]	0.641	0.630	[−0.0012, +0.0113]	0.003	5/5 (≤ 0.018)
1	[+0.062, +0.098]	0.630	0.627	[−0.0133, +0.0033]	0.021	4/5 (≤ 0.022)
2	[+0.060, +0.101]	0.633	0.627	[−0.0093, +0.0061]	0.023	5/5 (≤ 0.017)

sign: Fidelity on PhysioNet at +0.0113, HS-RBF on PhysioNet, 4 qubits at −0.0088 and Fidelity on Cho2017 at +0.0100. The first is one comparison of 35 that does not survive its family. The second deserves more than a mention, because it comes from the best-powered dataset in the study. On Cho2017, with 52 subjects, the parameter-free fidelity kernel leads the Riemannian-kernel twin by +0.0100 ($p = 0.001$, better in 39/52 subjects), while the other four kernels sit within −0.0008 to +0.0041 of it (smallest $p = 0.128$). This is the strongest evidence in the study of a quantum divergence separating from a metric-matched control in its own favour, and three facts bound it. It lies inside the ± 0.02 margin, so the two kernels are equivalent at ± 0.02 and distinguishable only because 52 subjects resolve a single accuracy point. It does not reach the best classical pipeline: on this dataset classical/TS+LR at 0.644 is itself ahead of the twin (−0.0061, $p = 0.010$), and the fidelity kernel’s lead over it, +0.0039 ($p = 0.174$), is not significant. And table 8 shows differences of this size changing sign between partitions of a smaller dataset. We record it, therefore, as a one-point advantage of the Bures geometry over the affine-invariant one on one dataset, which a practitioner choosing among SPD kernels may find useful, and not as evidence that quantum structure improves decoding.

Robustness to the partition. Every number above comes from one outer cross-validation partition. Table 8 repeats the PhysioNet three-qubit suite under two more. The frame effect reproduces (+0.060 to +0.102 over all 3 partitions), and so does the uncontrolled reversal: the best reference-frame kernel leads the best classical baseline by +0.0120 to +0.0177 in 3 of 3 partitions. The twin comparison, however, separates by kernel family rather than by partition, and it moves in both directions. The parameter-free kernels run −0.0023 to +0.0113 against the twin, the fidelity kernel leading it significantly in two of the 3 partitions (smallest $p = 0.003$), while the three bandwidth-parameterised kernels, which carry an extra hyperparameter to be tuned on 45 trials, fall behind by up to −0.0133 (smallest $p = 0.021$). Of the 15 comparisons across the three partitions, 3 are significant at 0.05: 2 in the quantum kernel’s favour and 1 in the twin’s. The equivalence bound for this dataset widens from 0.028 at the primary partition to 0.022 over all three.

We therefore state the PhysioNet result as a two-sided inequality rather than an equivalence, and the sign depends on which kernel is asked. A parameter-free quantum divergence is slightly ahead of the affine-invariant metric here, by about one accuracy point, and the same kernel is ahead by about the same amount on Cho2017, reported below; a bandwidth that must be tuned on 45 trials costs a kernel rather more than that against the same twin. Both effects are inside the ± 0.02 margin, neither changes which pipeline a practitioner should reach for, and we draw no conclusion about quantum structure from either: what separates fidelity from the Bures and relative-entropy kernels here is a hyperparameter, not a quantum resource.

We stress the counterfactual. Without this control, the same experiments support the claim that quantum kernels outperform tuned classical baselines at $p = 0.095$, on two datasets, with the gap widening on the better one. That claim would have been false. It is also, we suggest, the mechanism by which favourable results are obtained elsewhere in this literature: a quantum kernel that operates on whitened or aligned covariances, compared against a classical baseline that does not, will show an advantage of roughly the size reported in table 5 that has nothing to do with quantum structure.

3.6 Filter-bank baselines and larger registers

Since the reference-frame kernels reach the top of the table, the relevant question becomes whether they do so against the strongest classical family rather than against CSP and tangent-space alone. We therefore repeated the comparison with a filter bank of four sub-bands spanning the same passband, in the manner of filter-bank CSP [40], giving filter-bank CSP and filter-bank tangent-space baselines. The same construction enlarges the quantum register at no methodological




Figure 8. Accuracy in the sensor frame versus the reference frame. Thin lines are individual subjects, shown for the Fidelity kernel only. Every kernel improves on all three datasets; on IV-2a every kernel improves in every subject. The correction is worth two to three times more on IV-2a than on PhysioNet, and least on Cho2017, whose sensor-frame kernels start closest to the classical baselines.

Table 9. Cross-session transfer within subject on IV-2a. For each subject and in both directions, each model is trained on one recording session and tested on the other; the two directions are averaged per subject before any statistic is computed, so n is the number of subjects. Hyperparameters are chosen by stratified inner cross-validation inside the training session alone. In the reference frame each session is whitened by its own Fréchet mean, and for the test session that mean is computed from all of its trials. This uses test data, so it is stated here as well as in the text: it uses no labels, and it is exactly the unsupervised adaptation a deployed decoder can perform on the unlabelled data it receives before decoding anything. Δ is the gain from the reference frame, paired by subject and tested by Wilcoxon signed-rank.

Pipeline	Sensor	Reference	Δ	p	Better
QRE-RBF	0.576	0.760	+0.1844	0.004	9/9
Bures-RBF	0.573	0.760	+0.1867	0.004	9/9
Fidelity	0.542	0.759	+0.2172	0.004	9/9
riemann-kernel-SVM	0.734	0.755	+0.0204	0.012	8/9
logeuclid-kernel-SVM	0.724	0.755	+0.0309	0.020	8/9
HS-overlap	0.560	0.754	+0.1944	0.008	8/9
HS-RBF	0.578	0.746	+0.1686	0.004	9/9
TS+LR	0.717	0.738	+0.0216	0.039	8/9
MDM	0.701	0.732	+0.0312	0.023	7/9




Figure 9. The decisive control. Each quantum kernel is compared against a classical SPD-manifold kernel [39] used in the same support vector machine, on the same covariances, with the same tuning budget, in the same reference frame; only the metric differs. Points are mean paired differences with 90% confidence intervals, the interval that corresponds to the two one-sided tests at $\alpha = 0.05$, so this figure and table 7 count the same intervals as excluding zero. 3 of them exclude zero, drawn as orange diamonds, and every one lies inside the ± 0.02 margin (see text); table 8 repeats the PhysioNet block under two more partitions. Note the axis range: the entire plot spans eight accuracy points, against a frame effect (tables 5 and 9) of up to +0.217.

cost: four sub-bands over eight channels is a 32-dimensional covariance, hence a single density matrix on five qubits rather than three.

Parity survives. The best quantum kernel reaches 0.640 against 0.626 for filter-bank CSP, a difference of +0.0141 ($p = 0.014$), and the classical-twin control replicates at this register size as well (table 6). The frame effect is larger here than at three qubits, +0.110 to +0.128; in the sensor frame the five-qubit kernels fall to chance and are the only comparisons in that family to survive Holm correction, in the unfavourable direction.

The other route to a larger register is more channels. Running the same suite on sixteen sensorimotor channels, a four-qubit register, gives the same picture (tables 5 and 6): the frame is worth +0.066 to +0.103 ($n = 104$, all $p \leq 0.001$), the best reference-frame kernel, quantum/Bures-RBF-ref-SVM, reaches 0.632 against 0.634 for the twin, and the five differences span -0.0088 to -0.0017 (smallest $p = 0.048$). This register is also where the uncontrolled comparison is most seductive. Against classical/TS+LR the best reference-frame quantum kernel leads by +0.0071 ($p = 0.097$, better in 56/104 subjects), which unlike the three-qubit case (section 3.5) is conventionally significant; the classical twin leads the same baseline by +0.0088 ($p = 0.090$). Without the twin, four qubits would have supplied a significant quantum advantage on thirty subjects.

3.7 Cross-subject transfer

Table 10. Cross-subject transfer, leave-one-subject-out: each model is trained on the pooled trials of every other subject. Hyperparameters are selected by subject-grouped inner cross-validation on the training subjects only. In the reference frame each subject is whitened by its own Fréchet mean, which uses no labels and is therefore legitimate unsupervised adaptation for the held-out subject. Δ is the gain from the reference frame, paired by held-out subject. Every method improves; in the reference frame no method is distinguishable from any other.

Pipeline	Sensor	Reference	Δ	p	Better
HS-overlap	0.568	0.674	+0.1064	<0.001	80/104
TS+LR	0.626	0.674	+0.0472	<0.001	71/104
Bures-RBF	0.587	0.673	+0.0857	<0.001	77/104
HS-RBF	0.578	0.672	+0.0947	<0.001	78/104
Fidelity	0.579	0.672	+0.0925	<0.001	79/104
QRE-RBF	0.579	0.672	+0.0921	<0.001	75/104
logeuclid-kernel-SVM	0.617	0.669	+0.0526	<0.001	71/104
riemann-kernel-SVM	0.626	0.669	+0.0427	<0.001	70/104
MDM	0.546	0.662	+0.1162	<0.001	78/104

Within-subject accuracy is not the field's bottleneck; transfer to a new user is. It is also the setting in which a different geometry is most often conjectured to help. Table 10 and figure S3(a) report leave-one-subject-out evaluation over 104 subjects, training on the pooled trials of all others.

The invariance argument makes a sharp prediction here. Inter-subject shift is dominated by congruence, so in the sensor frame the quantum distances must be worse than the affine-invariant Riemannian one, and in the reference frame the congruence part of the shift can no longer separate them: parity is what the algebra predicts once the invariance is shared, and any residual difference must come from the part of the shift that is not a congruence, which the proposition does not cover. Both halves hold. Every method gains from recentring, the non-invariant quantum kernels gaining most; and in the reference frame the best method, quantum/HS-overlap at 0.674, is not distinguishable from any other. All nine geometries fall within 0.0118 of one another, with no comparison significant even before correction. This is a genuine null rather than an underpowered one: at $n = 104$ the two-sided Wilcoxon signed-rank test could have returned p as small as 9.9×10^{-32} .

3.8 Cross-session transfer

Cross-subject shift is dominated by congruence, so the parity just reported is what the algebra predicts. The complementary case is transfer across sessions within a subject: same head, same cap, similar impedances, so the congruence component is milder and what remains, drift in conditioning, effective rank and spectral shape, is exactly the residual that the invariance argument of section 2.6 says nothing about. It is where a quantum divergence, which weights the eigenvalue spectrum differently from the affine-invariant metric, would separate from its classical twin if it were going to anywhere, and we had named it in advance as the most promising place to look. IV-2a provides it: two recording sessions of 144 trials per subject. For each subject and in both directions we train on one session and test on the other (table 9, figure S2(a)); hyperparameters are tuned inside the training session only, and in the reference frame each session is whitened by its

own label-free mean, which for the test session is the unsupervised adaptation a deployed decoder can legitimately perform.

Two things happen. First, the frame matters more here than anywhere else in the study. In the sensor frame the quantum kernels reach only 0.542 to 0.578, close to chance, while the two classical baselines, tangent-space logistic regression and MDM, which do not see the session change as a change of frame, reach 0.701 to 0.717, and the two classical SPD kernels higher still (table 9). Recentring is worth +0.169 to +0.217 to the quantum kernels (all $p \leq 0.008$) and +0.022 to +0.031 to the classical ones. Session-to-session change evidently has a large congruence component after all, whatever its residual, and a kernel without the invariance is defenceless against it.

Second, once the frame is corrected the geometries tie again. The best pipeline is quantum/Bures-RBF at 0.760; the classical twin reaches 0.755 and the best conventional baseline, classical/TS+LR, 0.738. Against the twin the five quantum kernels differ by -0.0085 to +0.0054, and the smallest p among the five comparisons, $p = 0.500$, is the largest such value in any setting of the study; the five quantum kernels and the two classical twins lie within 0.0139 of one another. The equivalence bounds (table 7) are wider than in the within-subject IV-2a comparison, as they must be with one accuracy per subject rather than fifteen folds, but the largest is 0.0281, inside the study-wide ± 0.028 . The setting we identified in advance as the quantum divergences' best chance is therefore one more in which they match a classical kernel that differs only in metric. Either the non-congruence part of session shift is too small to matter at this register, or it is not something the Bures, fidelity or relative-entropy weightings exploit better than the affine-invariant one.

3.9 Four classes, not one binary contrast

Table 11. Four-class motor imagery on IV-2a (left hand, right hand, feet, tongue; chance = 0.25), $n = 9$ subjects, extended suite, protocol identical to the binary runs. Δ is the gain from the reference frame, paired by subject and tested by Wilcoxon signed-rank. Recentring is worth far more here than in the binary task, and the classical twin remains at least level with every quantum kernel afterwards, so neither conclusion of this paper depends on the task having been a single binary contrast.

Kernel	Sensor	Reference	Δ	p
Fidelity	0.346	0.633	+0.2875	0.004
HS overlap	0.357	0.626	+0.2687	0.004
HS-RBF	0.378	0.623	+0.2456	0.004
Bures-RBF	0.381	0.628	+0.2469	0.004
QRE-RBF	0.382	0.627	+0.2455	0.004
Riemannian kernel (twin)		0.636		
Log-Euclidean kernel		0.635		
TS + LR (best classical)		0.628		

Everything above decodes a single binary contrast, left hand against right, and a reader is entitled to ask whether the invariance argument is a property of that contrast rather than of the representation. We therefore repeated the extended suite on the four-class IV-2a task (left hand, right hand, feet, tongue; chance = 0.25), changing nothing but the label set.

The argument does not weaken with four classes, it sharpens. In the sensor frame the density-matrix kernels reach 0.346 to 0.382, which is a few points above chance and far below every tuned classical baseline: a kernel invariant only to rotations is close to useless on this task. Recentring is then worth +0.2455 to +0.2875 accuracy (all $p = 0.004$, every kernel improving in every one of the 9 subjects), against +0.064 to +0.102 for the same correction on the binary PhysioNet task. The frame is not a small correction here; it is most of the decoder.

And the control survives the change of paradigm. After recentring, the Riemannian twin reaches 0.636 and no quantum kernel passes it: the differences run -0.0123 to -0.0022, every one of them in the twin's favour, none significant (smallest $p = 0.098$). The twin also edges the best classical pipeline, TS + LR at 0.628. So a four-fold increase in the number of classes multiplies the value of the frame and leaves the quantum-versus-classical comparison exactly where the binary task left it. Neither of this paper's two conclusions depends on the task having been a single binary contrast.

3.10 Few-trial calibration: the one setting that separates

Every comparison so far estimates the reference state from a comfortable number of trials. Section 4.4 names the regime where that stops being true as the last place a quantum divergence could plausibly beat the affine-invariant metric: if the Fréchet mean is noisy, a divergence that weights a poorly estimated reference differently might degrade more gracefully. We stated that prediction before running the experiment and then ran it. For each subject a class-balanced test set is held out once, nested training sets of 10 to 160 trials are drawn ten times over, and every

Table 12. Few-trial calibration. For each subject a fixed, balanced test set is held out; each pipeline is tuned and trained on k balanced trials drawn from the rest (ten random draws, averaged per subject), so the reference state is estimated from those k trials alone. Columns give the Riemannian twin, tangent-space logistic regression, the best of the five reference-frame quantum kernels and the best sensor-frame one, the range of the five quantum-minus-twin differences, the smallest p among them, and the largest 90 % equivalence bound.

Dataset	k	Twin	TS+LR	Quantum	Sensor	Quantum – twin	min p	Bound
Cho2017 ($n = 52$)	10	0.566	0.563	0.575	0.560	[+0.0049, +0.0092]	0.003	0.0144
	20	0.585	0.584	0.593	0.573	[+0.0040, +0.0085]	0.018	0.0137
	40	0.596	0.610	0.612	0.589	[+0.0148, +0.0162]	<0.001	0.0249
	80	0.622	0.626	0.631	0.606	[−0.0004, +0.0084]	0.098	0.0179
	160	0.642	0.649	0.651	0.630	[−0.0130, +0.0090]	0.107	0.0287
IV-2a ($n = 9$)	10	0.671	0.661	0.685	0.536	[+0.0025, +0.0144]	0.035	0.0258
	20	0.689	0.688	0.702	0.553	[+0.0058, +0.0136]	0.055	0.0249
	40	0.714	0.705	0.718	0.552	[−0.0067, +0.0031]	0.344	0.0207
	80	0.736	0.730	0.738	0.581	[−0.0069, +0.0014]	0.520	0.0211
	160	0.745	0.749	0.749	0.585	[−0.0081, +0.0036]	0.195	0.0162

pipeline is tuned on each training set alone by the same inner search used everywhere else, with the reference state estimated from those k trials only.

The prediction holds, in the direction stated and at a size worth being precise about. On Cho2017 ($n = 52$), at the smallest calibration sets the five reference-frame kernels lead the twin by +0.0049 to +0.0092 (smallest $p = 0.003$), 0.575 in accuracy against the twin’s 0.566. By 160 trials the difference has gone and changed sign, spanning −0.0130 to +0.0090. Across the two datasets and every training size, 11 comparisons separate in the quantum kernel’s favour and 0 in the twin’s, the first asymmetry of that kind anywhere in this study.

Four things bound it, and they belong in the same breath as the result. The largest difference, 0.029, is inside the ± 0.02 margin at which the two families are equivalent everywhere else. It is a metric effect of about one accuracy point, not an advantage of a kind that would change practice. The kernels only clear the *best classical pipeline* at the smallest training set, and at the middling sizes where they beat the twin most clearly they merely match tangent-space logistic regression, which beats the twin there too. And the analysis we pre-specified, a per-subject slope of the difference against $\log_2 k$, is the weaker of the two: no kernel survives correction on it. We report the size-by-size comparison because it is better powered, not because it is more favourable, and we say so rather than presenting the stronger test as though it had been the plan. What the experiment settles is the paper’s own open question: the parity reported throughout is not universal, and the exception lies exactly where the geometry predicted it would.

It is worth noting that the same shape of result has just been reported in a domain with no connection to EEG. Benchmarking quantum and classical models for compound potency prediction across ten targets, Garg and Garg [46] find no overall separation between the best quantum and the best classical model (Wilcoxon $p = 0.189$), with quantum kernels holding a marginal edge confined to the sparse, data-poor targets. Two independent studies, different data, different quantum models, converging on parity overall and a small quantum edge only where training data is scarce, is a more useful signal about where these methods stand than either result alone.

3.11 Finite-shot estimation

All results so far assume noiseless, infinite-shot simulation. Since $\text{tr}(\rho\sigma)$ is the SWAP-test observable, S shots give an unbiased estimate with variance $(1 - k^2)/S$, which we simulate by sampling each unordered Gram entry binomially and projecting back to the positive semi-definite cone (table S4 and figure S3(b)). One departure from the main protocol should be stated: here the reference state is estimated from all of a subject’s trials rather than the training split alone. It is label-free, and the question is estimation cost rather than generalisation, so shot levels are comparable with one another; the absolute values are not comparable with the nested-CV benchmark and are not used that way.

The variance gain reported above does *not* translate into a proportionate saving in shots; the reference frame requires more shots to approach its own ceiling, because that ceiling is higher and the distinctions to be resolved are correspondingly finer. Normalising each frame by its own ceiling is in fact misleading, since the sensor frame converges quickly only because it converges to near-chance. The meaningful comparison is absolute: from 10^4 shots per Gram entry upwards, the reference frame exceeds what the sensor frame attains with unlimited shots. Below 10^3 shots both collapse to chance, so the frame raises the ceiling a well-resourced estimator can reach rather than rescuing a starved one.

Near-ceiling accuracy nonetheless requires 10^5 – 10^6 shots per Gram entry. A 45-trial subject has of order 10^3 unordered pairs, so a single Gram matrix costs 10^8 – 10^9 shots, and the IV-2a trial count raises that by a further factor of about forty. The reference frame makes these kernels work; it does not make them cheap.

3.12 Computational cost

The cost asymmetry survives the frame correction, which is the comparison that matters now that accuracy is at parity. In the sensor frame the gap is stark but partly beside the point, since the quantum kernels were also losing: classical/logvar+LDA reached its accuracy in 0.089 s per subject while quantum/IQP-kernel-SVM required 21.4 s, a factor of 240. In the reference frame the like-for-like comparison is against the classical twin, which delivers the same accuracy. On IV-2a the twin costs 15.7 s per subject; quantum/Bures-RBF-ref-SVM costs 117.7 s, a factor of 7.5, for an accuracy difference table 7 bounds at under ± 0.0109 . On PhysioNet the ratio is 3.8. The quantum kernel is not merely no better; it is the same answer at several times the price.

All of this is measured under noiseless, infinite-shot simulation. On hardware, estimating $\text{tr}(\rho\sigma)$ to precision ϵ costs $O(\epsilon^{-2})$ shots per Gram entry; section 3.11 puts the resulting budget at 10^5 – 10^6 shots per entry for near-ceiling accuracy, which the classical twin does not pay at all.

4 Discussion

4.1 Interpretation

Under matched tuning, matched features and paired statistics, quantum kernels did not improve motor-imagery decoding. That conclusion is consistent with the broader QML benchmarking literature [24], but the route by which we reach it differs from a bare negative result in a way we think matters, and it passes through a stage at which the opposite conclusion looked correct.

Evaluated in the sensor frame, the density-matrix kernels, the most principled construction here, requiring no dimensionality reduction and mapping onto a hardware-native SWAP test, performed *worst* of all pipelines, and section 3.3 appeared to explain why: EEG covariances are mutually so similar that their quantum overlaps saturate near unity. That explanation is incomplete. Concentration is a symptom of a deeper mismatch, and it is largely removable. Uhlmann fidelity, the Bures distance, the Hilbert–Schmidt overlap and the quantum relative entropy are invariant under *unitary* conjugation; EEG’s nuisances act by *congruence*. The affine-invariant Riemannian metric is invariant under that whole group, which is the structural reason tangent-space decoding is the classical state of the art. Comparing the two families in the sensor frame therefore measures, in part, a difference of invariance groups rather than a difference of geometries.

Correcting this (6) relieves concentration by 4.7–9.2 $\times$, lifts every kernel on all three datasets, and reverses the sign of the headline comparison. Taken alone, that is a publishable claim of quantum advantage at $p = 0.095$, replicated on a second dataset, with the effect *growing* on the better-powered one. We think it is worth being explicit that this is what our own experiments supported at that point, because the only thing standing between it and publication was a control that is not currently standard.

That control, an SPD kernel in the same classifier, same data, same frame, differing only in the metric, removes the claim entirely. Across seven settings the quantum kernels and their classical twins differ by at most 0.028, in both directions: under other partitions of PhysioNet the twin is significantly ahead of the bandwidth-tuned kernels, on Cho2017 the fidelity kernel is significantly ahead of the twin by one point, and neither excursion reaches the best classical pipeline. The twin is itself among the strongest pipelines we evaluated. The gain therefore belongs to the reference frame and to the kernel formulation, both of which are available classically and neither of which requires a quantum computer. Once the invariance groups are matched, the quantum geometries are not better; the parameter-free ones are not worse either, which is itself informative, since it locates the ceiling precisely where the algebra says it must be, and the bandwidth-tuned ones pay a small, partition-dependent price for their extra hyperparameter.

We draw two lessons. The first is specific: quantum information geometry is a mathematically exact description of EEG covariances that confers no decoding benefit over its classical counterparts, and reports to the contrary should be checked for a frame asymmetry before they are believed. The second is methodological: the dimension-matched control that this literature has begun to adopt is necessary but not sufficient. Its natural analogue for geometric methods, a *metric*-matched control, identical in every respect but the metric, is what distinguishes a claim about quantum structure from a claim about preprocessing.

4.2 Relation to prior work

Three strands of prior work bear on these results, and our relation to each is different.

The first is the theory of quantum kernels. Schuld [17] showed that supervised quantum models are kernel methods, so their merit reduces to whether the induced kernel suits the data; Huang *et al* [22] showed that any advantage is a property of the data as much as of the kernel; and Kübler *et al* [21] made the mechanism precise: a quantum kernel helps only where its inductive bias matches the target. Our results are an instance of that principle in which the inductive bias is an invariance group. The density-matrix kernels carry orthogonal invariance, EEG’s nuisances generate the congruence group, and the mismatch alone accounts for their sensor-frame deficit. Bandwidth tuning [36], which we apply throughout, and the exponential concentration of Thanasilp *et al* [25] address a different failure, one that grows with qubit count; table S3 shows it is not the failure at work here, since in the reference frame every density-matrix kernel has more Gram variance at six qubits than at three.

The second is the methodology of QML benchmarking. Bowles *et al* [24] found classical models matching or beating quantum ones across a broad suite and argued for ablations and matched baselines; we adopt both their entanglement ablation and the dimension-matched control. Our addition is the *metric*-matched control, which the geometric setting requires and the feature-map setting does not: a quantum kernel that operates on an SPD manifold has a classical counterpart on the same manifold, and comparing it with anything else conflates the metric with the representation.

The third is Riemannian BCI decoding [3, 5]. The reference-state correction is, algebraically, the re-centring of Zanini *et al* [27] and the alignment step of Riemannian Procrustes analysis [28], developed there as a transfer-learning heuristic. What we add is the reason it is not optional: for a family of kernels with the wrong invariance group it is the difference between a broken comparison and a fair one, and we prove that it restores affine invariance exactly rather than approximately. In that light our findings are compatible with reports of quantum benefit on EEG [6, 7] only under the interpretation that those benefits arise from the classical stages of the hybrid pipelines. Our dimension-matched control speaks to this directly, since a plain linear SVM on the same PCA-reduced tangent-space features matched or exceeded every circuit kernel, and our frame result names a second route to the same artefact: a quantum stage that operates on aligned covariances, compared with a classical baseline that does not. The pyRiemann-qiskit toolkit [26], which places a quantum classifier after tangent-space projection, is the natural place for the metric-matched control to be adopted.

4.3 Limitations

Five limitations bound our conclusions. (i) We study three datasets but one binary paradigm, left hand versus right hand; other motor-imagery classes and other BCI paradigms (P300, SSVEP) are untested. The trial-count concern that a single dataset would leave open is addressed by section 3.4. (ii) Our central claim is an equivalence, and equivalence is bounded rather than absolute. Two one-sided tests place the quantum kernels and their classical twins within ± 0.028 accuracy of one another across all 35 comparisons (table 7), so we can exclude differences larger than that, but not smaller ones; under two further partitions of PhysioNet the bound widens to ± 0.022 , every departure being in the twin’s favour (table 8). The TOST implementation is the paired t form and assumes the per-subject differences are approximately normal; at $n = 104$ that is mild, at $n = 9$ it is a real limitation and the IV-2a bounds should be read as indicative. (iii) Our invariance argument concerns congruence. Inter-subject and inter-session shift is dominated by, but not exhausted by, that group; differences in conditioning, effective rank and spectral shape survive recentring. Cross-session transfer, where that residual is proportionally largest, showed no quantum kernel separating from its twin (section 3.8), but it did so on one dataset of 9 subjects, which bounds the possibility rather than excluding it. (iv) The reference state is a Fréchet mean estimated from training data, so the correction is only as good as that estimate; its larger benefit on IV-2a (+0.164 to +0.187) than on PhysioNet (+0.064 to +0.102) is consistent with this, although Cho2017 (+0.012 to +0.063 from 200 trials) shows that how far the sensor frame sits from the reference frame matters as much as how well the reference is estimated. Either way the correction will be weaker in genuinely few-trial calibration settings. (v) Simulation is noiseless, and although section 3.11 adds shot noise it models sampling only, not gate error, decoherence or readout bias. All of these favour the quantum side, so none weakens a negative result; they do mean hardware performance would be worse, not better. We also survey a specific family of kernels: deep architectures such as EEGNet [47], and variational or quantum neural-network approaches, are outside our scope.

4.4 What would change the conclusion

Our claim is that the quantum geometries match their classical twins once the invariance groups agree, so what would overturn it is a setting in which a quantum divergence separates from a metric-matched classical comparator by more than the stated margin and reaches the best classical pipeline. Cho2017 supplies a separation of one accuracy point in the fidelity kernel’s favour (section 3.5) that does neither. We identified three candidates before running the extended experiments, and two of them have now been tried.

Non-congruence shift was the first, because it is exactly what our argument does not cover, and cross-session transfer is where it is proportionally largest: the congruence component is milder, so residual differences in conditioning and spectral shape, on which the quantum relative entropy and the Bures distance weight differently from the affine-invariant metric, carry more of the shift. Section 3.8 reports the outcome: parity, the closest quantum-versus-twin comparison reaching only $p = 0.500$. Larger registers were the third candidate, and the sweep of table S3 shows that the reference-frame kernels do not concentrate with qubit count. That removes a concern rather than the parity, since the metric-matched comparison at five qubits (section 3.6) had already tied.

What remains untried is the regime in which the reference state is poorly estimated. Our correction depends on a Fréchet mean, and the classical baselines depend on the same estimate; a divergence that degrades more gracefully under a noisy reference would be a real and practically relevant advantage, and few-trial calibration is where it would show. The gap between PhysioNet (+0.102 with 45 trials per subject) and IV-2a (+0.187 with 288), two datasets whose sensor-frame deficits are similar, is what a poorer reference-state estimate costs, and few-trial calibration is where that cost would bite hardest.

5 Conclusion

Treating EEG covariance matrices as quantum density matrices yields a principled, hardware-mappable family of quantum kernels that requires no dimensionality reduction. Evaluated in the sensor frame, these kernels are the worst pipelines in a controlled benchmark, and appear to be defeated by kernel concentration. That diagnosis is incomplete: the quantum quantities are invariant under unitary conjugation while EEG’s nuisances act by congruence, so the comparison was confounded by an invariance mismatch that the classical baselines do not share. Referring each state to a training-set Fréchet mean restores exact affine invariance, relieves concentration by 4.7–9.2 $\times$, improves every kernel on all three datasets, and reverses the sign of the classical-versus-quantum comparison.

It would be straightforward to stop there and report a quantum advantage. The control that prevents this, a classical SPD kernel identical in every respect but the metric, shows that across seven settings the quantum kernels stay within 0.028 of their classical twins in both directions: the bandwidth-tuned kernels behind under other partitions of PhysioNet, the parameter-free fidelity kernel about one accuracy point ahead on PhysioNet and again on the largest dataset, and none of them ever ahead of the best classical pipeline. Parity holds against filter-bank CSP, in cross-subject transfer, where all nine geometries fall within 0.0118, and in cross-session transfer, the setting we had named in advance as the most likely to show a real quantum effect. The improvement is real and useful, and it is entirely classical in origin.

We therefore find no evidence of quantum advantage for motor-imagery decoding, and we locate the ceiling: matched invariance permits parity and nothing more. For practitioners, the reference-state correction is worth adopting on its own terms. For the field, we suggest that a metric-matched control should accompany the dimension-matched one, since without it these same experiments support the opposite conclusion. All code and results are released to support replication.

Acknowledgments

The authors thank the PhysioNet, GigaDB and MOABB maintainers for curating and openly distributing the datasets used here.

Funding

No specific grant from any funding agency in the public, commercial or not-for-profit sectors was received for this work.

Author contributions

Md. Intiaj Alam Sajin: conceptualisation, data curation, formal analysis, investigation, methodology, resources, software, writing (original draft). **Md Wahiduzzaman Suva:**

investigation, resources, validation, writing (review and editing). **Esm E Moula Chowdhury Abha**: software, visualisation, writing (review and editing). **Samiul Islam Saif**: investigation, software, writing (review and editing). **Muhammad Hasibur Rashid Chayon**: supervision, writing (review and editing).

Data availability

The PhysioNet EEG Motor Movement/Imagery Database is publicly available at <https://physionet.org/content/eegmmidb/1.0.0/>, BCI Competition IV-2a at <https://www.bbci.de/competition/iv/>, and the Cho2017 dataset from GigaDB (<https://doi.org/10.5524/100295>); the latter two were accessed through MOABB. All analysis code, configuration and result files supporting this study, including the per-fold scores behind every table and figure, are openly available at <https://github.com/Imtiaj-Sajin/quantaEEG>.

Ethical statement

This study is a secondary analysis of three publicly available, de-identified EEG datasets [29, 30, 31, 33]. No new data were collected and the authors had no contact with any participant. The recordings were acquired by their original investigators under the ethical approval and informed-consent procedures described in the cited dataset publications, and were used here in accordance with the terms under which they are distributed.

Conflicts of interest

The authors declare no conflicts of interest.

References

- [1] Ramoser H, Müller-Gerking J and Pfurtscheller G 2000 *IEEE Transactions on Rehabilitation Engineering* **8** 441–446
- [2] Barachant A, Bonnet S, Congedo M and Jutten C 2012 *IEEE Transactions on Biomedical Engineering* **59** 920–928
- [3] Congedo M, Barachant A and Bhatia R 2017 *Brain-Computer Interfaces* **4** 155–174
- [4] Lotte F, Bougrain L, Cichocki A, Clerc M, Congedo M, Rakotomamonjy A and Yger F 2018 *Journal of Neural Engineering* **15** 031005
- [5] Yger F, Berar M and Lotte F 2017 *IEEE Transactions on Neural Systems and Rehabilitation Engineering* **25** 1753–1762
- [6] Chen C S, Chen S Y C, Tsai A H W and Wei C S 2024 153–158 (*Preprint* [2407.19214](https://arxiv.org/abs/2407.19214)) URL <https://arxiv.org/abs/2407.19214>
- [7] Olvera C, Ross O M and Rubio Y 2024 *Expert Systems with Applications* **247** 123354 URL <https://www.sciencedirect.com/science/article/abs/pii/S0957417424002197>
- [8] Lalawat R S, Bajaj V, Padhy P K and Lin C Y 2025 *Applied Acoustics* **240**
- [9] Kaur M, Kumar V and Upadhyay R 2026 *Soft Computing*
- [10] Koike-Akino T and Wang Y 2022 qEEGNet: Quantum AI for biosignal processing *2022 IEEE-EMBS International Conference on Biomedical and Health Informatics (BHI)*
- [11] Lins I D, Araújo M, Bezerra Souza Maior C, das Chagas Moura M, Ferreira-Martins A J, Chaves R and Canabarro A 2024 *Process Safety and Environmental Protection* **186** 1197–1213
- [12] Aksoy G, Cattani G, Chakraborty S and Karabatak M 2024 *Journal of Medical Systems* **48**
- [13] Carter R E, Wiczorek M A, Pacheco-Spann L M, Li F, Dougherty J M, Johnson P W, Wenzel B, Pandey M, Dalal A, Solano E, Yilmazer K, Soekadar S R, Thielen K R and Bruce C J 2026 *Mayo Clinic Proceedings*
- [14] Gupta R S, Wood C E, Engstrom T, Pole J D and Shrapnel S 2025 *npj Digital Medicine* **8** 237
- [15] Schuld M and Killoran N 2022 *PRX Quantum* **3** 030101
- [16] Schuld M and Killoran N 2019 *Physical Review Letters* **122** 040504

- [17] Schuld M 2021 *arXiv preprint (Preprint 2101.11020)*
- [18] McClean J R, Boixo S, Smelyanskiy V N, Babbush R and Neven H 2018 *Nature Communications* **9** 4812
- [19] Cerezo M, Arrasmith A, Babbush R, Benjamin S C, Endo S, Fujii K, McClean J R, Mitarai K, Yuan X, Cincio L and Coles P J 2021 *Nature Reviews Physics* **3** 625–644
- [20] Liu Y, Arunachalam S and Temme K 2021 *Nature Physics* **17** 1013–1017
- [21] Kübler J M, Buchholz S and Schölkopf B 2021 The inductive bias of quantum kernels *Advances in Neural Information Processing Systems* vol 34 pp 12661–12673
- [22] Huang H Y, Broughton M, Mohseni M, Babbush R, Boixo S, Neven H and McClean J R 2021 *Nature Communications* **12** 2631
- [23] Biamonte J, Wittek P, Pancotti N, Rebentrost P, Wiebe N and Lloyd S 2017 *Nature* **549** 195–202
- [24] Bowles J, Ahmed S and Schuld M 2024 *arXiv preprint (Preprint 2403.07059)* URL <https://arxiv.org/abs/2403.07059>
- [25] Thanasilp S, Wang S, Cerezo M and Holmes Z 2024 *Nature Communications* **15** 5200
- [26] Cattani G and Andreev A 2023 *Research Ideas and Outcomes* **9** e101006
- [27] Zanini P, Congedo M, Jutten C, Said S and Berthoumieu Y 2018 *IEEE Transactions on Biomedical Engineering* **65** 1107–1116
- [28] Rodrigues P L C, Jutten C and Congedo M 2019 *IEEE Transactions on Biomedical Engineering* **66** 2390–2401
- [29] Schalk G, McFarland D J, Hinterberger T, Birbaumer N and Wolpaw J R 2004 *IEEE Transactions on Biomedical Engineering* **51** 1034–1043
- [30] Goldberger A L, Amaral L A N, Glass L, Hausdorff J M, Ivanov P C, Mark R G, Mietus J E, Moody G B, Peng C K and Stanley H E 2000 *Circulation* **101** e215–e220
- [31] Tangermann M, Müller K R, Aertsen A, Birbaumer N, Braun C, Brunner C, Leeb R, Mehring C, Miller K J, Müller-Putz G R, Nolte G, Pfurtscheller G, Preissl H, Schalk G, Schlögl A, Vidaurre C, Waldert S and Blankertz B 2012 *Frontiers in Neuroscience* **6** 55
- [32] Jayaram V and Barachant A 2018 *Journal of Neural Engineering* **15** 066011
- [33] Cho H, Ahn M, Ahn S, Kwon M and Jun S C 2017 *GigaScience* **6** gix034
- [34] Uhlmann A 1976 *Reports on Mathematical Physics* **9** 273–279
- [35] Bhatia R, Jain T and Lim Y 2019 *Expositiones Mathematicae* **37** 165–191
- [36] Shaydulin R and Wild S M 2022 *Physical Review A* **106** 042407
- [37] He H and Wu D 2020 *IEEE Transactions on Biomedical Engineering* **67** 399–410
- [38] Havlíček V, Córcoles A D, Temme K, Harrow A W, Kandala A, Chow J M and Gambetta J M 2019 *Nature* **567** 209–212
- [39] Barachant A, Bonnet S, Congedo M and Jutten C 2013 *Neurocomputing* **112** 172–178
- [40] Ang K K, Chin Z Y, Zhang H and Guan C 2008 Filter bank common spatial pattern (FBCSP) in brain-computer interface 2008 *IEEE International Joint Conference on Neural Networks (IEEE World Congress on Computational Intelligence)* pp 2390–2397
- [41] Demšar J 2006 *Journal of Machine Learning Research* **7** 1–30
- [42] Holm S 1979 *Scandinavian Journal of Statistics* **6** 65–70
- [43] Gramfort A, Luessi M, Larson E, Engemann D A, Strohmeier D, Brodbeck C, Goj R, Jas M, Brooks T, Parkkonen L and Hämäläinen M 2013 *Frontiers in Neuroscience* **7** 267

- [44] Bergholm V, Izaac J, Schuld M, Gogolin C and Killoran N 2018 *arXiv preprint (Preprint [1811.04968](#))*
- [45] Chevallier S, Carrara I, Aristimunha B, Guetschel P, Sedlar S, Lopes B, Velut S, Khazem S and Moreau T 2024 *arXiv preprint (Preprint [2404.15319](#))*
- [46] Garg S and Garg S 2026 *ACS Medicinal Chemistry Letters* **17** 1736–1743
- [47] Lawhern V J, Solon A J, Waytowich N R, Gordon S M, Hung C P and Lance B J 2018 *Journal of Neural Engineering* **15** 056013