




Thin lines are individual subjects, shown for the Fidelity kernel only. Every kernel improves on every dataset, and on IV-2a every kernel improves in every subject. The correction is worth two to three times more on IV-2a than on PhysioNet, and least on Cho2017, whose sensor-frame kernels start closest to the classical baselines.